

Job Tracker Toolkit

Build your own automated job-board monitor in 30 minutes.
Free, durable, runs forever, costs \$0.

What you'll build

A daily automated job search that monitors company career pages for new postings matching your criteria, then emails you the matches the moment they appear. No subscriptions, no recruiter spam, no ads, just a clean filtered pipeline of opportunities at companies you actually care about.

What you need

- A computer with a web browser (any OS)
- A personal Gmail account
- ~25 minutes for first-time setup
- Access to Claude.ai (free works fine) for the intake step
- That's it. No coding required.

What you don't need

- Python or programming skill
- A code editor like Cursor or VS Code
- Command line / terminal experience
- Money. The whole stack runs on free tiers.

Created by Braden Tomlinson

[linkedin.com/in/bradentomlinson](https://www.linkedin.com/in/bradentomlinson)

MIT License — free to use, modify, and share with attribution

Contents

1.	What this is and who it's for	3
2.	The LLM intake — figure out what you're actually looking for	5
3.	Setting up your tools (GitHub, Gmail, secrets)	9
4.	Understanding the code (you don't have to read it)	14
5.	Configuring it for your search — let Claude do the work	17
6.	Adding more company sources	20
7.	Using your daily emails effectively	23
8.	Troubleshooting	25
9.	Limitations and what's next	27
10.	License and credits	29

1. What this is and who it's for

The Job Tracker Toolkit is a small, free tool that watches company job boards for new postings matching criteria you define. When new matches appear, you get an email with the title, location, salary (when available), and a direct link to apply.

It runs every day automatically on GitHub Actions. After the initial setup, you don't have to do anything to keep it running. You just open your inbox each morning and see what's new.

The problem it solves

Most job searches are reactive. You check LinkedIn occasionally, you remember to visit a couple of company pages when you have time, and by the time you find a great match the role has been open for three weeks and dozens of stronger candidates are already in interviews.

Senior-level roles especially move fast. Most fill within 2-3 weeks of posting. Being a same-day applicant is a real advantage. This toolkit makes that automatic.

Who this is for

- **Active job seekers** who want to stop missing relevant roles.
- **Passive job seekers** who only want to be alerted when their dream role at their dream company opens up.
- **Career changers** exploring a new industry or function.
- **People with 5-50 target companies** they want to track without manually checking each one.

Who this is NOT for

- People expecting AI to write applications for them. This tool tracks postings, not applies for you.
- People who want a guarantee. Tools find roles; you still have to compete for them.
- People unwilling to spend 25 minutes on setup. There's no shortcut.
- People searching exclusively at companies on unusual ATS systems (small or international companies often use bespoke job boards we don't currently support).

What you'll have at the end

- A private GitHub repository running your tracker daily.
- Email alerts when new matching postings appear at any company you've configured.
- Salary bands automatically extracted from postings when companies publish them.
- A system you can update anytime as your search criteria change.
- All of this for \$0 in monthly costs, indefinitely.

NOTE — ONE IMPORTANT CAVEAT

This tool currently supports three major job board systems: **Greenhouse**, **Ashby**, and **Workday**. Together these cover the majority of US-based tech, fintech, SaaS, healthcare, finance, and enterprise companies. If your dream company is on a less common system like Lever, iCIMS, SmartRecruiters, or Taleo, the toolkit can be extended (Section 6 explains how using Claude). But understand this isn't an out-of-the-box solution for every company on earth.

2. The LLM intake — figure out what you're actually looking for

Before you configure any code, you need clarity on what you're searching for. Vague criteria produce vague results. Most people skip this step and end up with a tracker that either floods them with irrelevant matches or stays silent for weeks.

The intake conversation below uses Claude (or ChatGPT, or any capable LLM) to translate your situation into a concrete tracker configuration: which companies, which keywords, which seniority level, which locations, which salary floor.

How to run the intake

Step 1. Open Claude at claude.ai in a new conversation. Free accounts work fine for this.

Step 2. Upload your most recent resume to the conversation. Drag-and-drop works.

Step 3. Paste the full prompt on the next page exactly as written. Don't edit it. The structure matters.

Step 4. Answer Claude's follow-up questions honestly. The whole conversation takes about 15-20 minutes. Don't rush it.

Step 5. Save the output. The final block will be a structured configuration you'll paste into the code in Section 5.

WHY THIS MATTERS

The output of this intake is the difference between a tracker that finds you 2-3 relevant roles per week and one that buries you under 30 irrelevant matches a day. The 20 minutes you invest here saves you hours every week.

The intake prompt (copy and paste this)

I'm setting up an automated job tracker that will monitor company job boards daily and email me when new matching roles are posted. Before I configure it, I need to figure out exactly what I'm looking for.

I've attached my resume. Please walk me through a structured intake conversation to clarify my job search criteria. Cover these topics, in this order, asking ONE topic at a time. Wait for my answer before moving to the next.

1. CURRENT SITUATION

- What's my current role, title, and approximate years of experience?
- Confirm what you see in my resume.
- Ask what I'm looking to change about my current situation.

2. CURRENT AND TARGET COMPENSATION

- What's my current total compensation (base + bonus + equity)?
- What's my realistic salary floor – the minimum I'd accept?
- What's my ideal salary target?
- Any non-salary deal-breakers? (Remote, equity, benefits, title, etc.)

3. TARGET ROLE TYPE

- What job titles am I targeting? (List 3-5.)
- What seniority level am I targeting?
- Am I open to lateral, step-up, or stretch roles?
- Any role types I want to specifically EXCLUDE?

4. TARGET INDUSTRY AND COMPANY STAGE

- What industries am I targeting? (e.g. fintech, healthcare, enterprise SaaS, edtech.)
- What company stage? (early startup / growth / late stage / public.)
- Any specific companies that are already on my list?
- Any companies I want to specifically AVOID?

5. LOCATION AND REMOTE

- Where do I live?
- Am I open to remote, hybrid, in-office?
- If hybrid or in-office, which cities am I open to?
- Am I open to relocation? Under what conditions?

6. WORK-LIFE PRIORITIES

- What's my ideal weekly hour commitment?
- Any specific cultural fit factors that matter? (e.g. flexible hours, no on-call, async-friendly, family-friendly.)

Once we've covered all six topics, produce a final structured output in this exact format:

JOB TRACKER CONFIGURATION

TARGET COMPANIES BY ATS SYSTEM:

Greenhouse: [list company slugs found at job-boards.greenhouse.io/<slug>]

Ashby: [list company slugs found at jobs.ashbyhq.com/<slug>]

Workday: [list companies – for each, format as: tenant/site/wd-prefix.

Example: [nvidia/NVIDIAExternalCareerSite/wd5](https://nvidia.com/NVIDIAExternalCareerSite/wd5)]

Other (not currently supported): [list these so I can add them later]

TITLE KEYWORDS (must contain ONE of these):

[list 5-10 lowercase words/phrases from target job titles]

SENIORITY KEYWORDS (must contain ONE of these):
[list of seniority terms appropriate to my target level]

LOCATION FILTERS (must contain ONE of these):
[list of location strings to allow]

TITLE EXCLUSIONS (skip if title contains any of these):
[list 5-10 lowercase words to filter out]

MINIMUM SALARY:
[my floor]

NOTES ON MY SEARCH:
[2-3 sentence summary of what I'm looking for and why, to remind
me later when I re-tune the tracker]

After producing the config, also tell me:

1. Three companies I should add to my list that I haven't mentioned yet, based on my background. For each, tell me what ATS they appear to use (you can guess based on company size and industry; I'll verify).
2. One concern about my search criteria that I might want to reconsider.

Begin with topic 1 now.

Example intake output

Here's a sample of what the intake output looks like, for a hypothetical product marketing manager looking to move from a mid-stage SaaS company to a senior IC or manager role at a growth-stage startup:

```
---
JOB TRACKER CONFIGURATION
---

TARGET COMPANIES BY ATS SYSTEM:
Greenhouse: notion, figma, linear, vercel, retool, airtable
Ashby: stripe, openai, anthropic
Workday: hubspot/HubSpotInc/wd1, atlassian/atlassian/wd103
Other (not currently supported):
  Webflow (uses Lever)
  Loom (acquired by Atlassian, see above)

TITLE KEYWORDS (must contain ONE of these):
product marketing, growth marketing, lifecycle marketing, content
marketing, demand generation, product launch, marketing manager

SENIORITY KEYWORDS (must contain ONE of these):
senior, lead, principal, staff, manager, head

LOCATION FILTERS (must contain ONE of these):
remote, united states, us, new york, san francisco

TITLE EXCLUSIONS (skip if title contains any of these):
intern, contractor, sales, engineer, designer, recruiter, customer
success, junior

MINIMUM SALARY:
$160,000

NOTES ON MY SEARCH:
Seven years in product marketing, currently at a mid-stage SaaS
company, looking for Senior IC or first-time Manager role at a
growth-stage startup. Remote-first preferred; would relocate to NYC
for the right role. Wants product-led companies, not sales-heavy
ones.
---
```

NEXT STEP — SAVE THIS OUTPUT

Copy the entire configuration block to a note somewhere you can find it. You'll paste pieces of it directly into the code in Section 5. The notes section at the bottom is for you, not the code — it helps you remember your intent when you update the tracker months later.

3. Setting up your tools

Three setup steps, in order: create your GitHub account, get a Gmail app password, then upload the toolkit files. Budget 25 minutes total. Don't skip steps. The order matters.

IMPORTANT — PRIVACY AND SECURITY

Keep this completely separate from your current employer. Use your personal Gmail. Create a brand-new personal GitHub account if you don't have one. Do not use a work computer for setup if your employer monitors employee devices. Once set up, the toolkit runs in the cloud — you never have to touch it again from any device.

3.1 — Create your GitHub account (5 min)

GitHub is where your tracker code will live, and where the daily automation runs (for free, forever). If you already have a personal GitHub account separate from work, skip to 3.2.

1. Go to <https://github.com/signup>.
2. Use your **personal Gmail**. Not work email.
3. Create a strong password and pick a username. Keep it professional — recruiters may see it.
4. Verify your email when prompted.
5. When asked about plan, choose **Free**.
6. Skip or briefly answer the interests questions. They don't affect anything.

3.2 — Get a Gmail App Password (5 min)

Gmail requires a separate app password to allow automated email sending. Your normal Gmail password will not work for this.

1. Go to <https://myaccount.google.com/security>.
2. If **2-Step Verification** is OFF, turn it ON. Required before app passwords are available. Follow Google's prompts.
3. Go to <https://myaccount.google.com/apppasswords>.
4. In **App name**, type: **Job Tracker**
5. Click **Create**.
6. Google shows you a 16-character password. **Copy it now** — you cannot view it again later.

SECURITY — TREAT THIS PASSWORD LIKE A CREDENTIAL

This app password gives any system holding it the ability to send email from your Gmail account. Don't paste it into chat apps, documents, or anywhere public. We'll put it directly into GitHub's secure secrets storage in the next step, where it's encrypted and never visible again.

3.3 — Create your tracker repository (10 min)

Now we'll create a private repository on GitHub and add the three files that make up the toolkit.

Create the repo

1. Go to <https://github.com/new> (must be logged in).
2. Repository name: **job-tracker**
3. Description: **My daily job tracker**
4. Set visibility to **Private**. Important — your search criteria and email address should not be public.
5. Check the box **Add a README file**.
6. Click **Create repository**.

Add the `job_tracker.py` file

1. On your new repo page, click **Add file** → **Create new file**.
2. In the filename box, type: **job_tracker.py**
3. Paste the entire contents of the `job_tracker.py` file (in Section 4 of this guide, or download from the repo URL in Section 10).
4. Scroll down. Click **Commit changes** twice.

Add the workflow file (the daily scheduler)

This file goes in a subfolder. GitHub will create the subfolder automatically when you type the slashes in the filename.

1. Click **Add file** → **Create new file**.
2. In the filename box, type exactly: **.github/workflows/track.yml**
3. Paste the workflow YAML (also in Section 4).
4. Click **Commit changes** twice.

Verify the structure

Your repo should now look like this:

```
job-tracker/
├── .github/
│   └── workflows/
│       └── track.yml
├── README.md (you can leave this alone or delete it)
└── job_tracker.py
```

3.4 — Add your Gmail password as a GitHub secret (3 min)

1. On your repo page, click **Settings** (top-right of repo).
2. In the left sidebar, click **Secrets and variables** → **Actions**.
3. Click **New repository secret** (green button).
4. **Name** (exactly): **GMAIL_APP_PASSWORD**
5. **Secret**: paste your 16-character Gmail app password. Spaces don't matter.
6. Click **Add secret**.

You should now see GMAIL_APP_PASSWORD listed under Repository secrets. The value is encrypted — you can't view it after this, only delete or update it.

3.5 — Run your first scan (2 min)

1. Click **Actions** at the top of your repo.
2. If GitHub shows a banner asking to enable workflows, click the green button to enable.
3. In the left sidebar, click **Daily Job Tracker**.
4. On the right, click **Run workflow** → green **Run workflow** button.
5. Wait ~30 seconds. Refresh. A run will appear.
6. Click into the run, then click **track** to expand the steps. Click **Run job tracker** to see output.

Expected result on first run: the tracker will run, but find zero matches because you haven't configured any companies yet. That's fine. Section 5 covers configuration.

SETUP COMPLETE

If your first run succeeded (green checkmark), the infrastructure is working. From here, the tracker runs automatically every day at 7am Mountain Time. You only ever need to come back to update your configuration as your search evolves.

4. Understanding the code (you don't have to read it)

This section contains the two files that make up the toolkit. You don't need to read or understand the code to use it — Section 5 explains exactly which lines you edit, and how.

If you want to make changes beyond what Section 5 covers, you can paste the code into a new Claude conversation and ask for help. The code is documented heavily for that exact purpose.

File 1 — `job_tracker.py` (the main script)

About 430 lines. Defines the four main pieces: fetchers (Greenhouse, Ashby, Workday), the filter system, salary extraction, and the email-sending logic. The configuration block lives at the top of the file.

Get the full file from the toolkit's GitHub repository (URL in Section 10) and paste it into your `job_tracker.py` file.

File 2 — `.github/workflows/track.yml`

About 40 lines. Tells GitHub Actions to run the script daily at 7am Mountain Time (14:00 UTC). To change the time, edit the cron schedule near the top.

```
name: Daily Job Tracker

on:
  schedule:
    - cron: "0 14 * * *"
  workflow_dispatch: {}

permissions:
  contents: write

jobs:
  track:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout repository
        uses: actions/checkout@v4

      - name: Set up Python
        uses: actions/setup-python@v5
        with:
          python-version: "3.11"

      - name: Run job tracker
        env:
          GMAIL_APP_PASSWORD: ${ secrets.GMAIL_APP_PASSWORD }
        run: python job_tracker.py

      - name: Commit updated state file
        run: |
          git config --local user.email "actions@github.com"
          git config --local user.name "Job Tracker Bot"
          git add seen_jobs.json
          if git diff --staged --quiet; then
            echo "No state changes to commit."
          else
```

```
    git commit -m "Update tracker state ($(date -u +%Y-%m-%d))"
    git push
fi
```

What each tool does

- **GitHub Actions:** runs the script every morning. Free for personal use.
- **Python script:** fetches jobs from each company, filters them, and sends emails.
- **seen_jobs.json:** a small state file that remembers which jobs you've already been notified about. Updated automatically every run.
- **Gmail app password:** lets the script send you email. Stored encrypted as a GitHub secret.

5. Configuring it for your search (let Claude do the work)

This is where most setup guides ask you to edit code. We're not going to do that. Instead, you'll let Claude turn your intake output (from Section 2) into a fully configured `job_tracker.py` file. Your job is to copy and paste. That's it.

WHY THIS APPROACH

Editing Python by hand means watching for matching quotes, commas, brackets, and indentation. One wrong comma and the whole thing breaks with an error message you can't read. Claude doesn't make those mistakes. Let it do the syntax. You stay focused on what your search should actually look for.

Step 1 — Grab the current `job_tracker.py` file

Before Claude can update the file, it needs to see what's in it. Get the current content from your repo:

1. Go to your repo on github.com.
2. Click `job_tracker.py` to open it.
3. Click the **copy** icon (top-right of the file viewer — looks like two overlapping squares). This copies the entire file to your clipboard.

Step 2 — Open a fresh Claude conversation

Go to claude.ai and start a new chat. Don't reuse the conversation you did the intake in — start fresh so Claude isn't distracted by earlier discussion.

Step 3 — Paste this prompt into Claude

Copy the entire prompt below exactly. Wherever you see square brackets like **[paste here]**, replace the whole thing (including the brackets) with what they say.

I have a job tracker script (`job_tracker.py`) that needs to be updated with my search configuration. Below I'll give you:

1. My intake configuration output (what to filter for)
2. My personal Gmail address (for email notifications)
3. The current contents of my `job_tracker.py` file

Your job is to update `job_tracker.py` with my search configuration, keeping everything else in the file exactly the same. Then return the complete updated file as a single code block I can copy and paste back into GitHub.

Specifically, you should update these things in the CONFIG block:

- `GREENHOUSE_COMPANIES`: the Greenhouse companies from my intake
- `ASHBY_COMPANIES`: the Ashby companies from my intake
- `WORKDAY_COMPANIES`: the Workday companies from my intake, properly formatted as dictionaries with `tenant`, `site`, and `wd` fields
- `KEYWORDS`: the title keywords from my intake
- `SENIORITY`: the seniority keywords from my intake
- `LOCATIONS_OK`: the location filters from my intake
- `TITLE_EXCLUDE`: the title exclusions from my intake
- `ENABLE_EMAIL`: change to `True`
- `EMAIL_FROM` and `EMAIL_TO`: set to my Gmail address

Do not change any other code. Do not change function definitions, fetchers, filters, or main logic. Only update the CONFIG block values at the top.

Return the entire updated file as one code block, with no commentary before or after. I will copy the entire block and paste it into GitHub.

=== MY INTAKE CONFIGURATION ===

[paste your entire JOB TRACKER CONFIGURATION block from Section 2 here, including all the company lists, keywords, seniority terms, locations, exclusions, and salary minimum]

=== MY PERSONAL GMAIL ADDRESS ===

[paste your personal Gmail address, for example: jane.smith@gmail.com]

=== CURRENT job_tracker.py FILE ===

[paste the entire contents of your job_tracker.py file that you copied from GitHub in Step 1]

Step 4 — Claude returns your updated file

Claude will respond with a single code block containing the entire updated job_tracker.py file. The block starts with the import statements at the top and ends with the bottom of the file.

Click the **copy** button in the top-right corner of Claude's code block to copy the entire file to your clipboard.

SANITY CHECK

Before pasting back into GitHub, glance at Claude's response. The code block should be hundreds of lines long — the entire file. If you only see the CONFIG block, ask Claude: *"Please return the full file, not just the CONFIG block."* If the file looks dramatically different from what you copied in Step 1, ask Claude what it changed.

Step 5 — Paste the updated file into GitHub

1. Go back to **job_tracker.py** on github.com.
2. Click the pencil icon (top-right) to edit.
3. Select all the existing content (Ctrl+A or Cmd+A) and delete it.
4. Paste your clipboard (the updated file from Claude).
5. Scroll down. Click **Commit changes** twice.

Step 6 — Test it

1. Click **Actions** at the top of your repo.
2. Click **Daily Job Tracker** in the left sidebar.
3. Click **Run workflow** → green **Run workflow** button.
4. Wait ~30 seconds. Check your email within a couple minutes.

Expected first run: the email lists every currently-matching job because the tracker has no history yet. From the next run onward, only NEW jobs are emailed.

If something doesn't work

Problem: Claude returns just the CONFIG block, not the full file.

Fix: Reply to Claude with: *"Please return the entire updated file from import statements through the end. I need to paste the whole thing back into GitHub."*

Problem: The workflow failed when you ran it.

Fix: Click into the failed run, find the red X step, click it to expand the error. Then go back to Claude (same conversation) and paste: *"My workflow failed with this error: [paste error]. What's wrong with the file you generated?"* Claude will identify the issue and give you a corrected file.

Problem: The run succeeded but no email arrived.

Fix: Open the Actions log and look for **[WARN] GMAIL_APP_PASSWORD not set**. If you see it, go to Settings → Secrets and variables → Actions and confirm the secret is named exactly **GMAIL_APP_PASSWORD**. Also check your Gmail spam folder. If no warning appears, also confirm Claude set **ENABLE_EMAIL = True** in your updated file.

Updating your search later

Your search criteria will change as you learn more about what's out there. To update the tracker any time:

1. Re-run the LLM intake from Section 2 (or just edit your saved configuration).
2. Repeat Steps 1-5 of this section with the new configuration.
3. Commit. The next run picks up your changes.

Doing this once a month keeps your tracker tuned to your actual search. The minute the criteria stop matching reality, the daily emails stop being useful.

WORTH REPEATING — YOU NEVER HAVE TO EDIT CODE

Every update to your tracker — adding companies, changing keywords, switching from full-time to part-time targets, narrowing by salary — goes through Claude. You describe what you want in plain English. Claude returns an updated file. You paste it back. If you never look at the Python code itself, the tool still works perfectly for you.

6. Adding more company sources

As your search evolves, you'll find new companies you want to track. Adding them is a smaller version of the same Claude-driven flow from Section 5.

The general workflow

1. Identify the company and its job board URL.
2. Open a fresh Claude conversation. Paste the prompt below, replacing the brackets with your details.
3. Paste Claude's response into `job_tracker.py` on GitHub.
4. Commit. The next run picks it up.

The universal "add a company" prompt

This single prompt handles any company on any ATS. Claude figures out which list to add it to, formats it correctly, and returns the updated file.

I want to add a new company to my job tracker. Here's what I know:

Company name: [company name, e.g. Salesforce]
Their career page URL: [the URL where you see their job listings,
e.g. https://salesforce.wd12.myworkdayjobs.com/External_Career_Site]

Please:

1. Tell me which ATS this company uses (Greenhouse, Ashby, Workday, or something else).
2. If it's Greenhouse, Ashby, or Workday, add it to my `job_tracker.py` in the correct list with the correct format. Return the entire updated file as a single code block.
3. If it's a different ATS (Lever, iCIMS, SmartRecruiters, etc.), tell me clearly that the toolkit doesn't yet support it, and ask if I want you to add support for that new ATS – but don't do it automatically.

Here's my current `job_tracker.py` file:

[paste your current `job_tracker.py` file here – grab it from GitHub the same way you did in Section 5, Step 1]

How to find the right URL

Visit the company's careers or jobs page. Click any job listing to open it. The URL of the specific job (not the careers homepage) is what Claude needs. Copy and paste that URL into the prompt.

TIP — COMPANIES OFTEN HAVE MULTIPLE CAREER SITES

A few large companies post jobs on multiple sites — an internal Workday instance and a public Greenhouse board, for example. If your search keeps missing roles you know exist, ask Claude: *"Does [company name] post jobs on more than one job board? If so, what are all their URLs?"* You may need to add multiple entries for one company.

Adding support for a new ATS

If a company you really want to track uses an ATS the toolkit doesn't yet support (Lever, iCIMS, SmartRecruiters, Taleo, BambooHR, etc.), Claude can extend the toolkit for you. Use this prompt:

```
I want to add support for a new ATS to my job tracker. Right now
it supports Greenhouse, Ashby, and Workday. I want to add: [Lever /
iCIMS / SmartRecruiters / Taleo / etc.]
```

```
A company that uses this ATS is [company name], and their public
careers URL is [paste URL].
```

Please:

1. Research the public API or job listing format used by this ATS.
2. Write a new fetcher function called `fetch_[atsname]` that matches the pattern of `fetch_greenhouse`, `fetch_ashby`, and `fetch_workday` in my existing file.
3. Add a new CONFIG list called `[ATSNAME]_COMPANIES` at the top of the file.
4. Add a new loop in the `main()` function to call `fetch_[atsname]` for each company in that list.
5. Return the entire updated `job_tracker.py` file as a single code block.

Here's my current `job_tracker.py` file:

[paste your current `job_tracker.py` file here]

Once Claude returns the updated file, paste it back into GitHub (Section 5, Step 5). Then you can add companies on the new ATS using the universal "add a company" prompt above.

NOTE — A CAVEAT ON NEW ATS SUPPORT

Some ATS systems don't expose public APIs and require web scraping, which is slower and more fragile. Claude will tell you if that's the case. For ATSs without clean APIs, you may get occasional run failures that need debugging. Greenhouse, Ashby, and Workday are the most reliable.

7. Using your daily emails effectively

The tracker is only as useful as your response to its emails. Here's how to get the most from it.

Sample daily email

Found 2 new matching job(s) on 2026-05-11:

- [NOTION] Senior Product Marketing Manager (GREENHOUSE)
 - Remote within United States
 - \$170,000-\$210,000
 - <https://job-boards.greenhouse.io/notion/jobs/...>
- [STRIPE] Lifecycle Marketing Lead, Growth (ASHBY)
 - San Francisco, CA or Remote within United States
 - \$185,000-\$240,000
 - <https://jobs.ashbyhq.com/stripe/...>

The 24-hour rule

When a real match arrives, apply within 24 hours. Senior-level roles especially fill within 2-3 weeks of posting. Being a first-week applicant materially improves your odds.

If the role is clearly a fit, apply same-day. Have your resume ready, draft a quick tailored cover letter (Claude can help in 10 minutes), and submit. Don't perfect — submit.

Track your applications separately

The toolkit tracks *new postings*, not *your applications*. Use a simple spreadsheet, Notion table, or Airtable to track what you've applied to. Columns to include:

- Company
- Role title
- Date applied
- Application status
- Recruiter / contact name
- Next-action date
- Salary band (if disclosed)

Refresh your config every 30 days

Your search criteria evolve as you learn what's actually out there. Once a month, re-run the LLM intake from Section 2 with your updated thinking and refresh your CONFIG block. Add companies you've discovered, remove ones that aren't hiring, adjust keywords based on what's actually getting matched.

When to retire the tracker

When you accept an offer, you have two options: pause the workflow (Settings → Actions → disable workflow) or leave it running. I recommend leaving it running for 6-12 months after you start your new role. It gives you ambient awareness of the market and confirms whether you made the right move. Reactivating later if you want to job search again costs nothing.

8. Troubleshooting

Workflow failed (red X in Actions)

Click the failed run, find the step with the red X, copy the error message, and paste it to Claude with this prompt:

```
My GitHub Actions job tracker workflow failed. Here's the error:
```

```
[paste the error message]
```

```
What should I check or change?
```

No emails arriving

- Check the Actions log. If you see **[WARN] GMAIL_APP_PASSWORD not set**, the secret name is wrong. Re-do Section 3.4 with the exact name `GMAIL_APP_PASSWORD`.
- Confirm `ENABLE_EMAIL = True` at the top of `job_tracker.py`.
- Check your spam folder.
- Make sure both `EMAIL_FROM` and `EMAIL_TO` are set to your real Gmail.

Too many matches every day

- Add more words to `TITLE_EXCLUDE`.
- Make `KEYWORDS` more specific — instead of 'manager', try 'product manager' or 'marketing manager'.
- Tighten `LOCATIONS_OK` — remove broad terms like 'us' that match everything.

Zero matches for a week

- `KEYWORDS` may be too narrow. Add synonyms.
- `LOCATIONS_OK` may exclude too much. Try setting it to `[]` temporarily to see what comes back.
- Some companies just aren't hiring right now. Check the volume of total jobs in the Actions log.

A company returns HTTP 404 every run

- The slug is wrong. Visit the company's career page and verify the exact slug from the URL.
- For Greenhouse: confirm the URL pattern is exactly **`job-boards.greenhouse.io/SLUG`**. Some companies use the older `boards.greenhouse.io` subdomain, which can have slightly different behavior.
- For Workday: double-check the `wd` prefix. Workday URLs use `wd1`, `wd5`, `wd103`, `wd2pr`, and other variations.

How to ask Claude for help with any other problem

When something's not working and you can't tell what, open a new Claude conversation and paste this:

```
I have a Python job tracker on GitHub Actions. Here's what's happening:
```

```
[describe the problem in plain English]
```

```
Here's the Actions log error (if any):
```

[paste the error or note that there's no error]

Here's the relevant part of my `job_tracker.py`:

[paste the CONFIG block, plus any fetcher function involved]

What should I change?

9. Limitations and what's next

What this toolkit doesn't do

- **Doesn't auto-apply.** Auto-applying creates spam, harms your reputation, and produces zero offers. Application quality matters.
- **Doesn't track which roles you've applied to.** Use a separate spreadsheet or Notion table.
- **Doesn't know when applications close.** Roles can stay on job boards after they're effectively filled.
- **Doesn't capture every ATS.** Greenhouse, Ashby, and Workday cover most major US tech, fintech, and enterprise companies, but Lever, iCIMS, SmartRecruiters, Taleo, BambooHR, and bespoke career sites are not covered out of the box. Section 6 explains how to extend.
- **Doesn't tailor your resume.** Use Claude or ChatGPT separately for that.
- **Doesn't write cover letters.** Same — use Claude separately.
- **Doesn't reach out to recruiters.** That's your job. The tool tells you when something opens; you take the action.

Reasonable extensions you can build

Once you've used this for a few weeks and want more, here are natural next steps you can build with Claude's help:

- **Slack notifications** instead of (or alongside) email.
- **Daily digest at a different time** — change the cron schedule in `track.yml`.
- **Per-company keyword filters** — track different keywords for different companies.
- **A weekly summary email** on Friday rolling up the week's matches.
- **Auto-tagging by salary band** — high-confidence vs unknown comp.
- **Integration with a Notion or Airtable database** as a CRM for your search.
- **Adding more ATS systems** — Lever, iCIMS, SmartRecruiters.

Things to be careful about

- **Don't make the repo public** while your config contains specific salary numbers, email addresses, or company targets that you don't want to share.
- **Don't commit your Gmail password into the code.** Always use GitHub Secrets.
- **Don't run the tracker more than once per day.** GitHub Actions has free tier limits, and aggressive polling makes you look like a bot to the company's career site.
- **Don't share the `seen_jobs.json` file** — it doesn't contain sensitive data per se, but it's specific to your search history.

10. License and credits

License

This toolkit is released under the **MIT License**. In plain English:

- You can use it for personal job searching, free, forever.
- You can modify it, share it, and redistribute it.
- You can use it commercially (e.g., build a paid service on top of it) as long as the original license and credit remain in the code.
- You can't hold the author liable if something goes wrong.

Attribution

If you share this toolkit, please credit the creator. A line like “Built on the Job Tracker Toolkit by Braden Tomlinson” or a link to the LinkedIn post that introduced it is sufficient.

About the creator

Braden Tomlinson is a credit risk operations manager at a fintech in Las Vegas. He built the original version of this toolkit for his own job search and decided to share it after realizing how many people were doing job searches without leverage. Connect on LinkedIn at [linkedin.com/in/bradentomlinson](https://www.linkedin.com/in/bradentomlinson).

Get the latest version

The latest version of the code, plus update notes, will be shared from his LinkedIn profile. If you find this useful, share it with someone else who's job-searching.

Final note

Tools don't get you offers. Work does. This toolkit removes the friction of finding opportunities so you can spend your time on the things that actually matter: applying with care, networking with intent, and preparing for interviews. Use the time you save wisely.

Good luck.